

Exploratory Data Analysis of Anscombe's Quartet

Author: Avveer Singh Khurana

Date: 2025-10-21

GitHub: [AvveerSchoolAccount](#)

Notebook / Code: [GitHub Repository](#)

Short Description

This notebook performs an exploratory data analysis (EDA) on **Anscombe's Quartet** — a dataset that demonstrates how identical summary statistics can hide very different data patterns.

We compute statistics, plot visualizations, and discuss why **EDA must include visual inspection**.

Abstract / Executive Summary

Anscombe's Quartet (Francis Anscombe, 1973) shows that four datasets can share identical means, variances, correlations, and regression lines — yet look completely different when plotted.

This notebook demonstrates the importance of visualization in exploratory data analysis (EDA) by computing these statistics and visualizing patterns using scatter, residual, and distribution plots.

Introduction

Anscombe's Quartet was created by statistician *Francis Anscombe* in 1973 to highlight the importance of graphing data before analyzing it.

Each of the four datasets has nearly identical summary statistics (mean, variance, correlation, regression line, and R^2), but the data points form completely different patterns when plotted.

The purpose of this analysis is to:

- Compute key summary statistics for each dataset
- Visualize the datasets using multiple plot types
- Show that **statistics alone can be misleading without visualization**

Data

The dataset used here is **Anscombe's Quartet**, consisting of four small datasets labeled I, II, III, and IV.

Each dataset contains pairs of x and y values.

For reproducibility, the dataset was recreated directly within the notebook instead of being loaded from a file.

```
In [1]: # --- Imports ---
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm

# --- Step 1: Create the dataset ---
data = {
    'x123': [10, 8, 13, 9, 11, 14, 6, 4, 12, 7, 5],
    'y1': [8.04, 6.95, 7.58, 8.81, 8.33, 9.96, 7.24, 4.26, 10.84, 4.82, 5],
    'y2': [9.14, 8.14, 8.74, 8.77, 9.26, 8.1, 6.13, 3.1, 9.13, 7.26, 4.74],
    'y3': [7.46, 6.77, 12.74, 7.11, 7.81, 8.84, 6.08, 5.39, 8.15, 6.42, 5],
    'x4': [8, 8, 8, 8, 8, 8, 19, 8, 8, 8],
    'y4': [6.58, 5.76, 7.71, 8.84, 8.47, 7.04, 5.25, 12.5, 5.56, 7.91, No]
}

df = pd.DataFrame(data)
df.head()
```

```
Out[1]:   x123    y1    y2    y3    x4    y4
0     10  8.04  9.14  7.46     8  6.58
1      8  6.95  8.14  6.77     8  5.76
2     13  7.58  8.74 12.74     8  7.71
3      9  8.81  8.77  7.11     8  8.84
4     11  8.33  9.26  7.81     8  8.47
```

Methods

We calculate the following summary statistics for each dataset:

- Mean
- Variance
- Standard deviation
- Covariance
- Pearson correlation coefficient
- Linear regression slope, intercept, and R^2

Summary Statistics Formulas

Mean: $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$

Sample Variance: $s_x^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$

Sample Standard Deviation: [$s_x = \sqrt{s_x^2}$]

Covariance: [$\text{cov}(x, y) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$]

Pearson Correlation: [$r_{xy} = \frac{\text{cov}(x, y)}{s_x s_y}$]

Linear Regression (Least Squares): [$y = b_0 + b_1 x$] where

[$b_1 = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$, $b_0 = \bar{y} - b_1 \bar{x}$]

Coefficient of Determination (R^2): [$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$]

```
In [2]: # --- Step 2: Reshape the data (melt) ---
df_melted = pd.melt(
    df,
    id_vars=['x123', 'x4'],
    value_vars=['y1', 'y2', 'y3', 'y4'],
    var_name='dataset',
    value_name='y'
)

# Assign correct x column for each dataset
df_melted['x'] = df_melted.apply(
    lambda row: row['x123'] if row['dataset'] in ['y1', 'y2', 'y3'] else
    axis=1
)

df_melted.drop(columns=['x123', 'x4'], inplace=True)

# --- Step 3: Define summary statistics function ---
def summary_stats(g):
    n = g.shape[0]
    mean_x = g['x'].mean()
    mean_y = g['y'].mean()
    var_x = g['x'].var(ddof=1)
    var_y = g['y'].var(ddof=1)
    std_x = g['x'].std(ddof=1)
    std_y = g['y'].std(ddof=1)
    cov = g[['x', 'y']].cov().iloc[0, 1]
    corr = g['x'].corr(g['y'])

    X = sm.add_constant(g['x'])
    model = sm.OLS(g['y'], X).fit()
    slope = model.params['x']
    intercept = model.params['const']
    r2 = model.rsquared

    return pd.Series({
        'n': n,
        'mean_x': mean_x, 'mean_y': mean_y,
        'var_x': var_x, 'var_y': var_y,
        'std_x': std_x, 'std_y': std_y,
        'cov_xy': cov, 'r_xy': corr,
        'slope': slope, 'intercept': intercept, 'r2': r2
    })

# --- Step 4: Compute summary table ---
```

```
summary_table = df_melted.groupby('dataset').apply(lambda g: summary_stats(g))
summary_table
```

/tmp/ipykernel_13640/960195626.py:46: FutureWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, and in a future version of pandas the grouping columns will be excluded from the operation. Either pass `include_groups=False` to exclude the groupings or explicitly select the grouping columns after groupby to silence this warning.

```
summary_table = df_melted.groupby('dataset').apply(lambda g: summary_stats(g)).round(4)
```

Out[2]:

| | n | mean_x | mean_y | var_x | var_y | std_x | std_y | cov_xy | r_xy | slope |
|---------|------|--------|--------|-------|--------|--------|--------|--------|--------|--------|
| dataset | | | | | | | | | | |
| y1 | 11.0 | 9.0 | 7.5009 | 11.0 | 4.1273 | 3.3166 | 2.0316 | 5.5010 | 0.8164 | 0.5001 |
| y2 | 11.0 | 9.0 | 7.5009 | 11.0 | 4.1276 | 3.3166 | 2.0317 | 5.5000 | 0.8162 | 0.5000 |
| y3 | 11.0 | 9.0 | 7.5000 | 11.0 | 4.1226 | 3.3166 | 2.0304 | 5.4970 | 0.8163 | 0.4997 |
| y4 | 11.0 | 9.0 | 7.5620 | 11.0 | 4.5358 | 3.3166 | 2.1297 | 6.0353 | 0.8147 | NaN |

Visualizations

Below are plots that reveal how different the datasets truly are, despite having identical summary statistics.

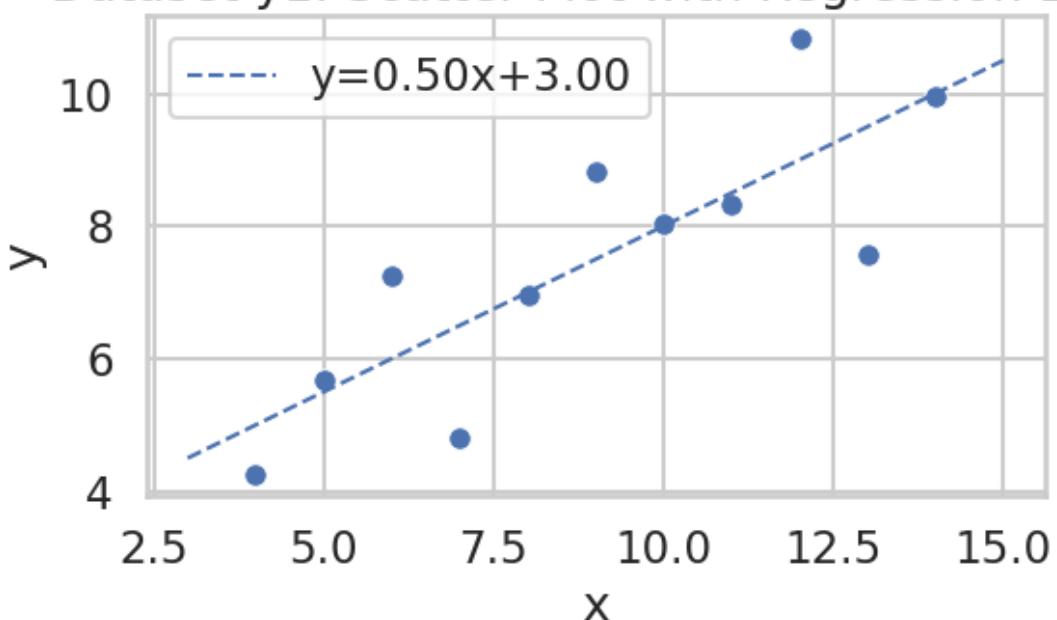
1. Scatter plots with regression lines
2. Residual plots
3. Overlaid scatter plot for comparison
4. Violin plots showing Y distribution across datasets

In [3]:

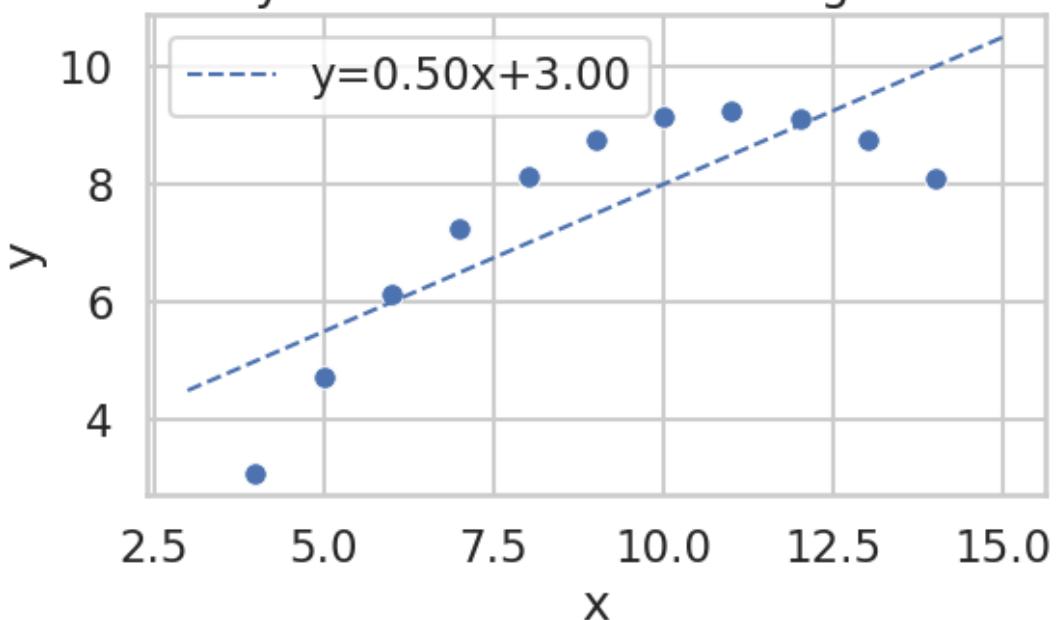
```
sns.set(style="whitegrid", context="talk")

# Scatter plots with regression line for each dataset
for name, group in df_melted.groupby('dataset'):
    plt.figure(figsize=(6, 4))
    sns.scatterplot(data=group, x='x', y='y', s=70)
    X = sm.add_constant(group['x'])
    model = sm.OLS(group['y'], X).fit()
    xs = np.linspace(group['x'].min() - 1, group['x'].max() + 1, 100)
    ys = model.params['const'] + model.params['x'] * xs
    plt.plot(xs, ys, '--', linewidth=1.5, label=f"y={model.params['x']:.2f}")
    plt.title(f"Dataset {name}: Scatter Plot with Regression Line")
    plt.xlabel("x")
    plt.ylabel("y")
    plt.legend()
    plt.tight_layout()
    plt.show()
```

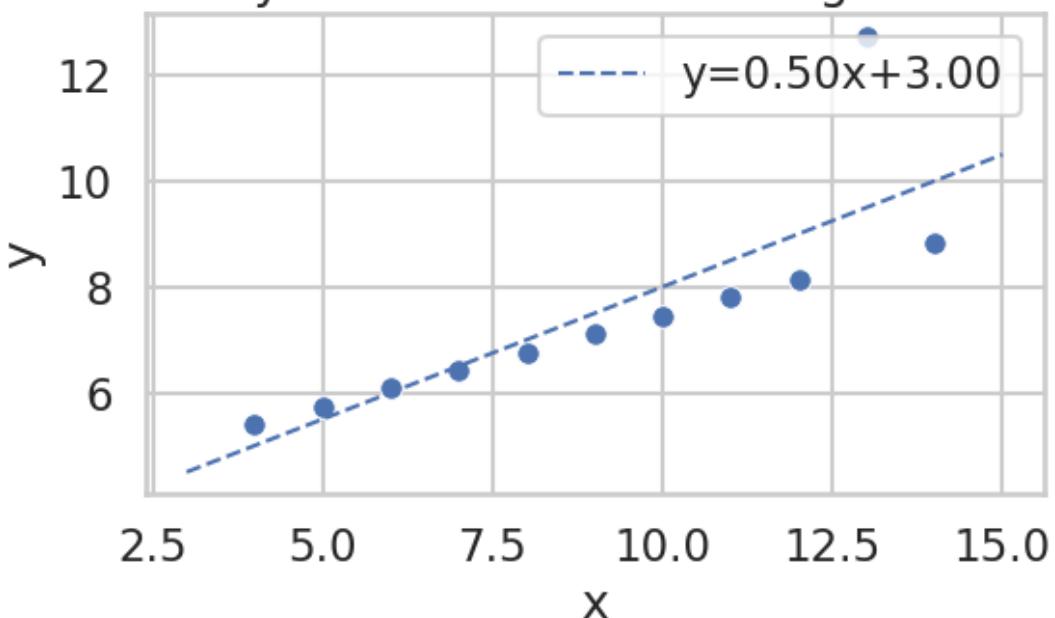
Dataset y1: Scatter Plot with Regression Line



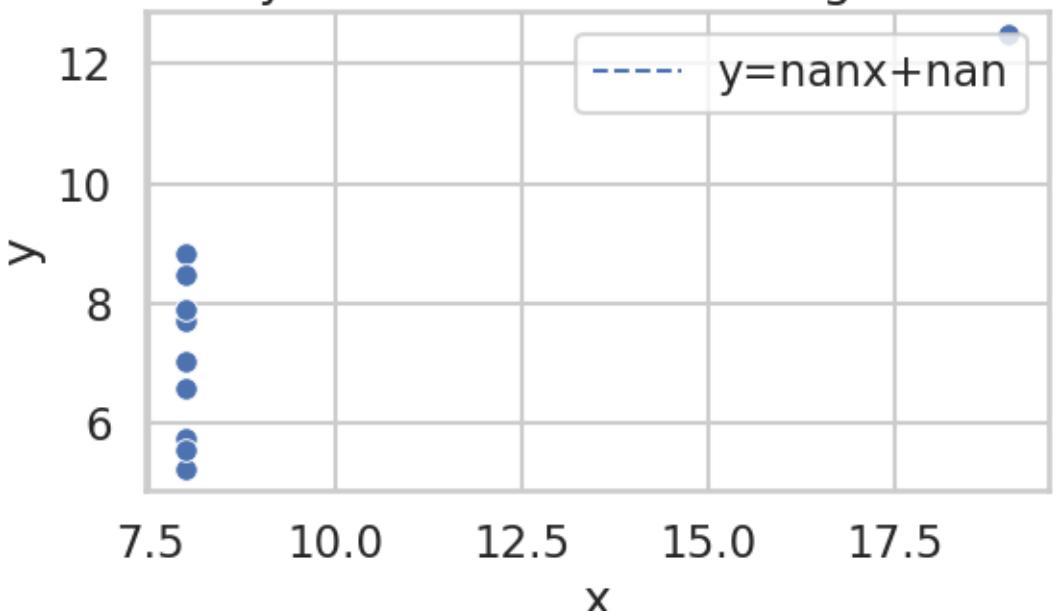
Dataset y2: Scatter Plot with Regression Line



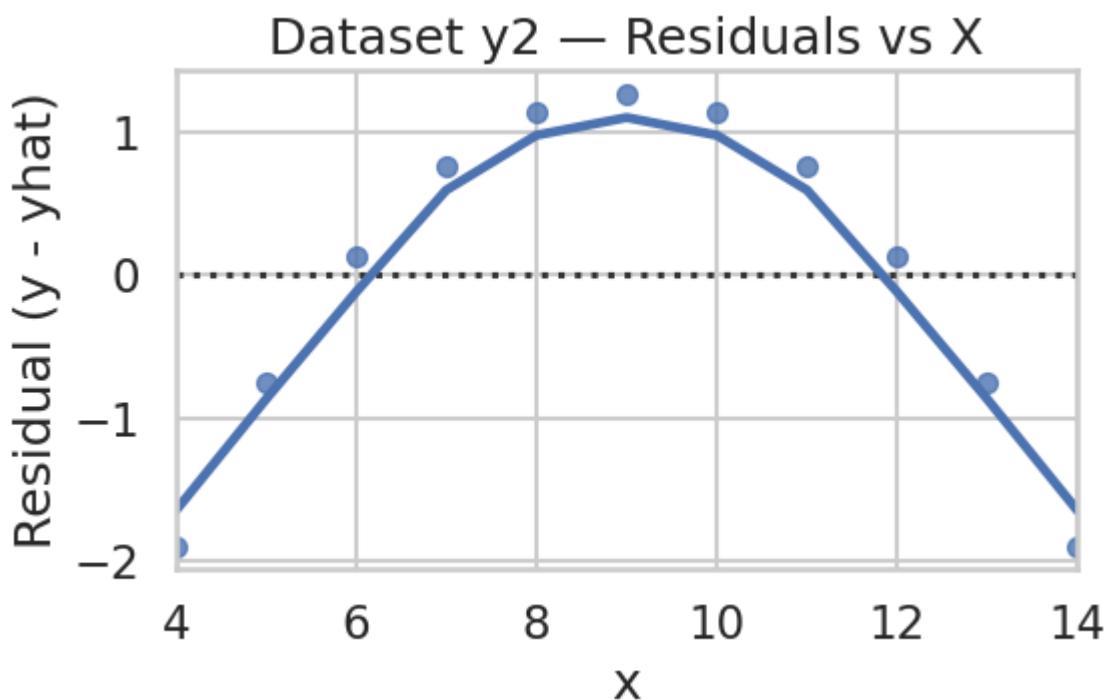
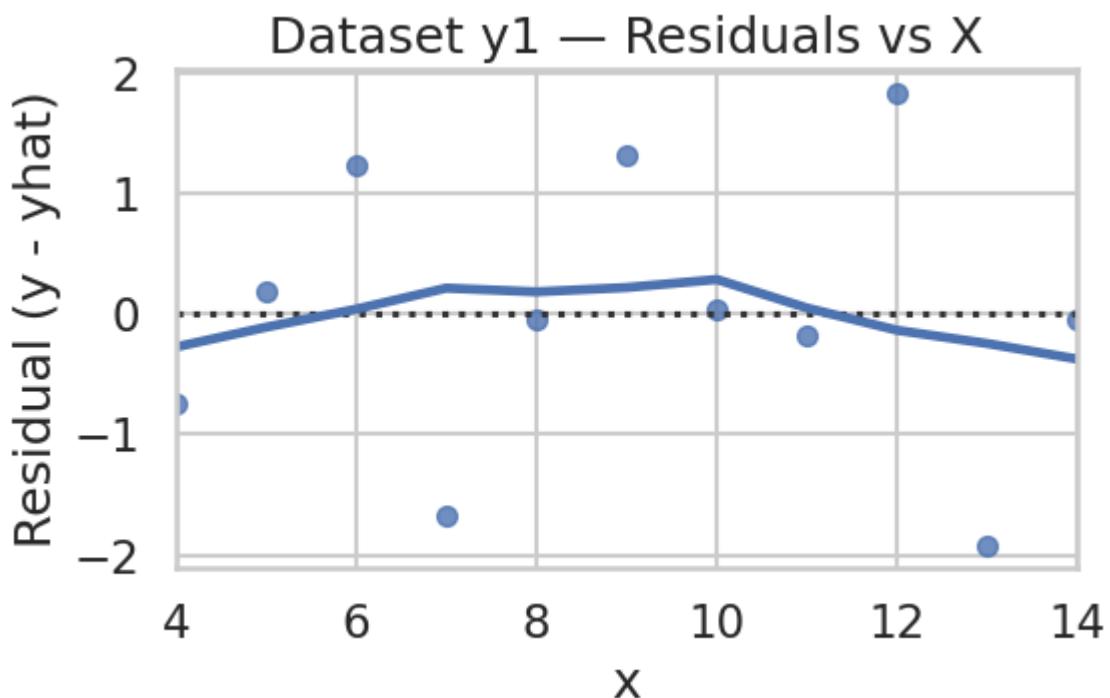
Dataset y3: Scatter Plot with Regression Line

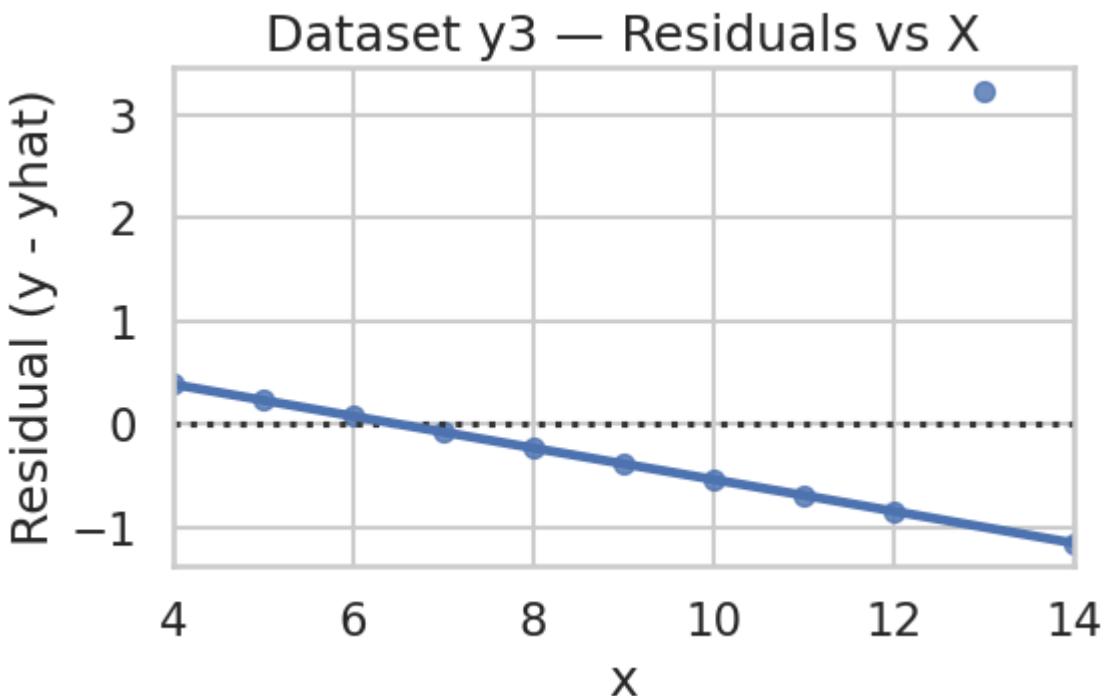


Dataset y4: Scatter Plot with Regression Line

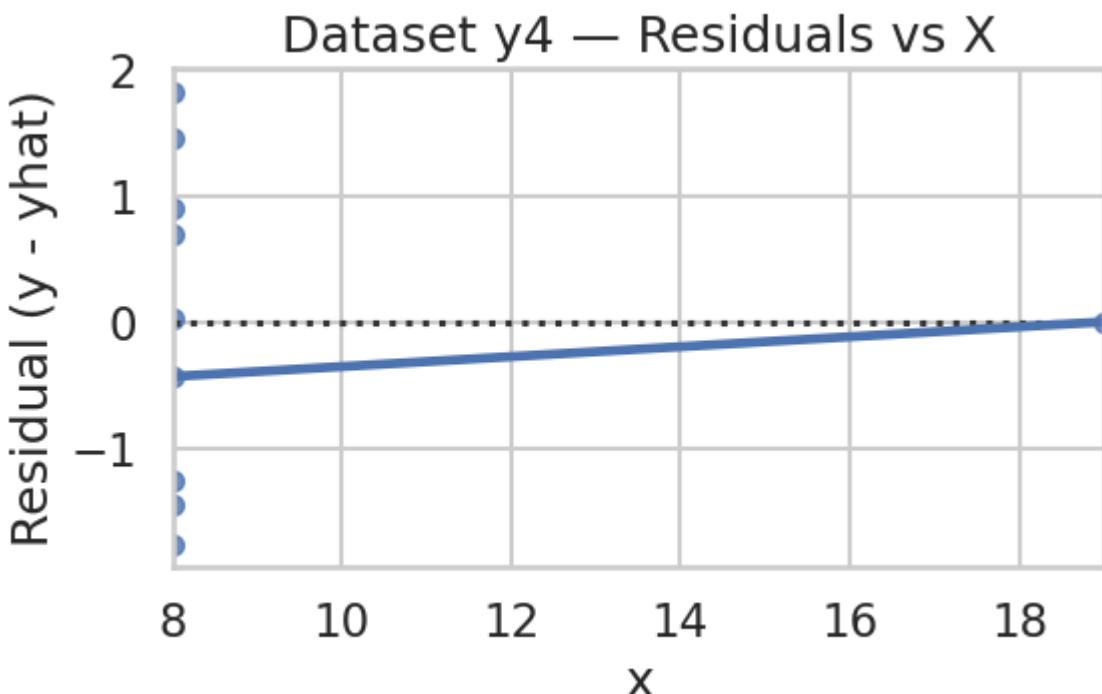


```
In [4]: # Residual plots
for name, group in df_melted.groupby('dataset'):
    X = sm.add_constant(group['x'])
    model = sm.OLS(group['y'], X).fit()
    plt.figure(figsize=(6, 4))
    sns.residplot(x='x', y='y', data=group, lowess=True, scatter_kws={'s': 100})
    plt.title(f"Dataset {name} - Residuals vs X")
    plt.xlabel("x")
    plt.ylabel("Residual (y - yhat)")
    plt.tight_layout()
    plt.show()
```





```
/home/avveer-singh-khurana/anscombe-project/venv/lib/python3.12/site-packages/statsmodels/nonparametric/smoothers_lowess.py:226: RuntimeWarning: invalid value encountered in divide  
    res, _ = _lowess(y, x, x, np.ones_like(x),
```

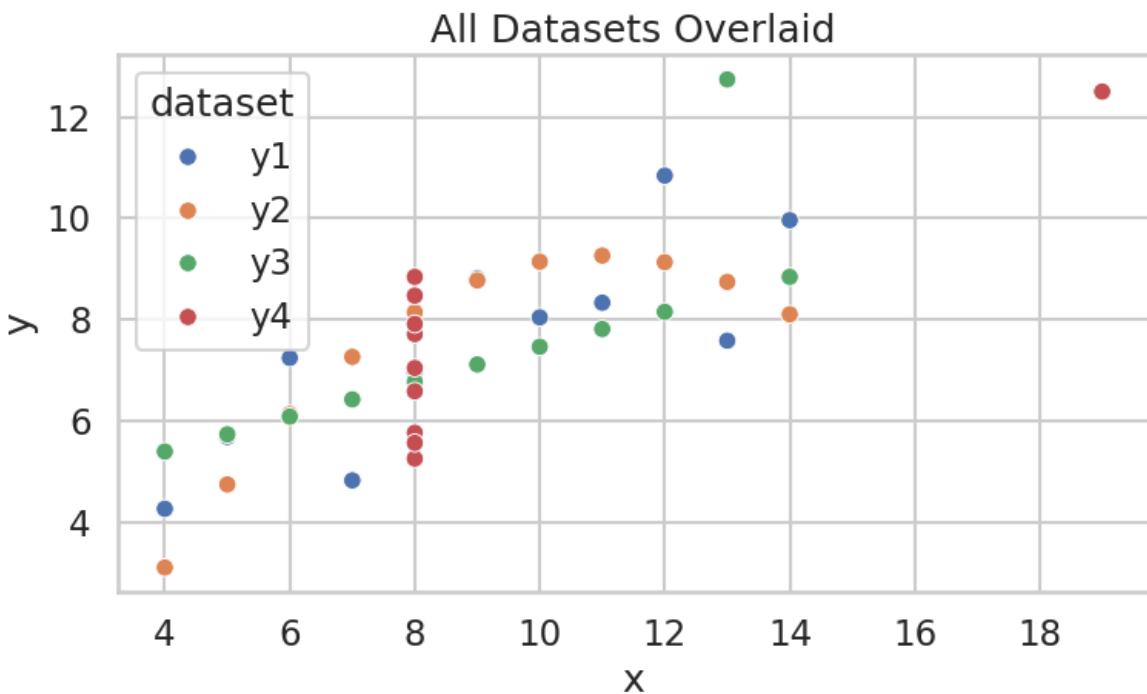


```
In [5]: # Overlaid scatter comparison  
plt.figure(figsize=(8, 5))  
sns.scatterplot(data=df_melted, x='x', y='y', hue='dataset', s=70)  
plt.title("All Datasets Overlaid")  
plt.xlabel("x")  
plt.ylabel("y")  
plt.tight_layout()  
plt.show()  
  
# Violin plot of Y distributions  
plt.figure(figsize=(8, 5))
```

```

sns.violinplot(data=df_melted, x='dataset', y='y', palette='muted')
plt.title("Distribution of Y Across Datasets")
plt.xlabel("Dataset")
plt.ylabel("Y Value")
plt.tight_layout()
plt.show()

```



/tmp/ipykernel_13640/3771097717.py:12: FutureWarning:

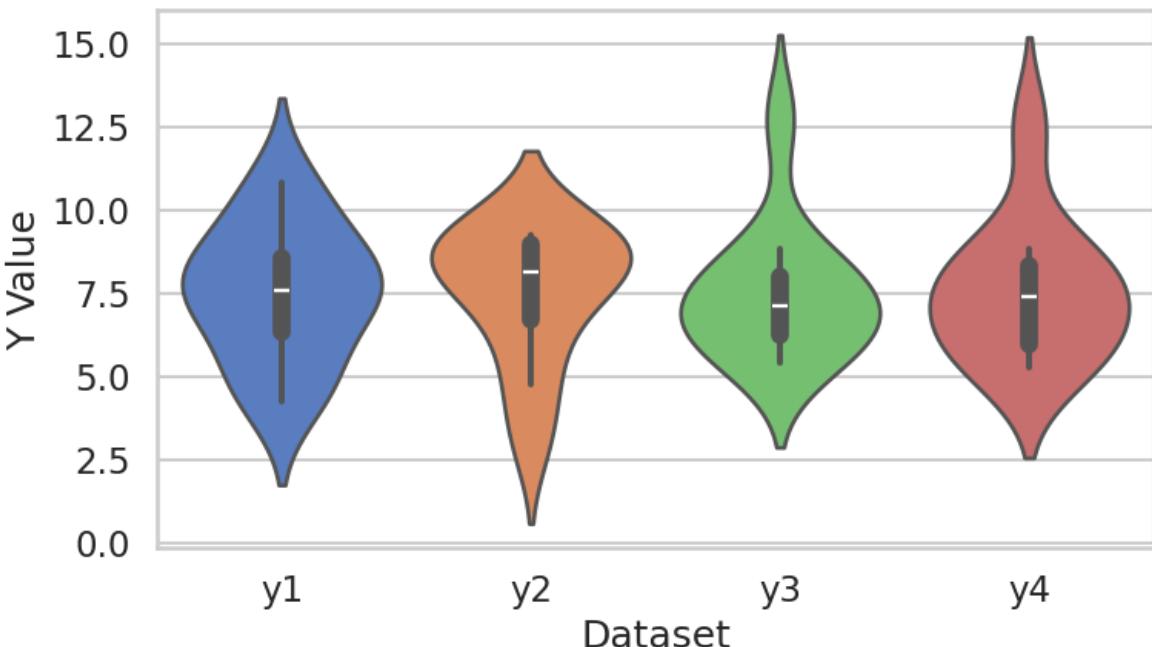
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.violinplot(data=df_melted, x='dataset', y='y', palette='muted')

```

Distribution of Y Across Datasets



Interpretation

The summary statistics of all four datasets are nearly identical:

- Same mean and variance for x and y
- Same correlation (≈ 0.816)
- Same linear regression line and R^2

However, when plotted:

- Dataset I shows a normal linear relationship.
- Dataset II is clearly **nonlinear**.
- Dataset III has a single **outlier** that changes the line.
- Dataset IV has all points identical except one **extreme outlier**.

This proves that **summary statistics alone are insufficient** — visualization is essential for accurate data understanding.

Reproducibility & Code

The full notebook and source code are available here:

[🔗 GitHub Repository](#)

This notebook was created in **Python** using:

- pandas, numpy
- seaborn, matplotlib
- statsmodels

Collaboration Notes

This project was completed individually by **Avveer Singh Khurana**.

All code, analysis, and writing are original.

Appendix

Full code is included in this notebook.

For online access, visit the GitHub repository below:

[🔗 https://github.com/AvveerSchoolAccount/anscombe-eda](#)

All figures and plots were generated automatically from the Python code in this file.

In []: